



# Creating Databases and Tables



- We've focused on querying and reading data from existing databases and tables.
- Let's now shift our focus to creating our own databases and tables.



- Section Overview
  - Data Types
  - Primary and Foreign Keys
  - Constraints
  - CREATE
  - INSERT
  - UPDATE
  - DELETE, ALTER, DROP



- We first focus on learning a few theoretical concepts, such as choosing the correct data type for a stored value and setting possible constraints on it.
- We will also learn about primary and foreign keys.



# Data Types



- We've already encountered a variety of data types, let's quickly review the main data types in SQL.



- Boolean
  - True or False
- Character
  - char, varchar, and text
- Numeric
  - integer and floating-point number
- Temporal
  - date, time, timestamp, and interval



- UUID
  - Universally Unique Identifiers
- Array
  - Stores an array of strings, numbers, etc.
- JSON
- Hstore key-value pair
- Special types such as network address and geometric data.





- When creating databases and tables, you should carefully consider which data types should be used for the data to be stored.
- Review the documentation to see limitations of data types:
- **[postgresql.org/docs/current/datatype.htm](https://www.postgresql.org/docs/current/datatype.htm)**  
|



- For example
  - Imagine we want to store a phone number, should it be stored as numeric?
  - If so, which type of numeric?
- We could take a look at the documentation for options...



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



- Based on the limitations, you may think it makes sense to store it as a **BIGINT** data type, but we should really be thinking what is best for the situation.
- Why bother with numerics at all?
- We don't perform arithmetic with numbers, so it probably makes more sense as a **VARCHAR** data type instead.



- In fact, searching for best practice online, you will discover its usually recommended to store as a text based data type due to a variety of issues
  - No arithmetic performed
  - Leading zeros could cause issues, 7 and 07 treated same numerically, but are not the same phone number



- When creating a database and table, take your time to plan for long term storage
- Remember you can always remove historical information you've decided you aren't using, but you can't go back in time to add in information!



# Primary and Foreign Keys



- A primary key is a column or a group of columns used to identify a row uniquely in a table.
- For example, in our dvdrental database we saw customers had a unique, non-null `customer_id` column as their primary key.





- Primary keys are also important since they allow us to easily discern what columns should be used for joining tables together.



- Example of Primary Key

Query Editor

Query History

1

SELECT \* FROM customer

Data Output

Explain

Messages

Notifications

	customer_id [PK] integer	store_id smallint	first_name character varying (45)	last_name character varying (45)
1	524	1	Jared	Ely
2	1	1	Mary	Smith
3	2	1	Patricia	Johnson
4	3	1	Linda	Williams



- Example of Primary Key

Query Editor

Query History

1

SELECT \* FROM customer

Data Output

Explain

Messages

Notifications

	customer_id [PK] integer	store_id smallint	first_name character varying (45)	last_name character varying (45)
1	524	1	Jared	Ely
2	1	1	Mary	Smith
3	2	1	Patricia	Johnson
4	3	1	Linda	Williams



- Notice its integer based and unique

Query Editor

Query History

1

SELECT \* FROM customer

Data Output

Explain

Messages

Notifications

	customer_id [PK] integer	store_id smallint	first_name character varying (45)	last_name character varying (45)
1	524	1	Jared	Ely
2	1	1	Mary	Smith
3	2	1	Patricia	Johnson
4	3	1	Linda	Williams



- Later we will learn about SERIAL data type

Query Editor

Query History

1

SELECT \* FROM customer

Data Output

Explain

Messages

Notifications

	customer_id [PK] integer	store_id smallint	first_name character varying (45)	last_name character varying (45)
1	524	1	Jared	Ely
2	1	1	Mary	Smith
3	2	1	Patricia	Johnson
4	3	1	Linda	Williams



- A foreign key is a field or group of fields in a table that uniquely identifies a row in another table.
- A foreign key is defined in a table that references to the primary key of the other table.



- The table that contains the foreign key is called referencing table or child table.
- The table to which the foreign key references is called referenced table or parent table.
- A table can have multiple foreign keys depending on its relationships with other tables.



- Recall in the dvdrental database payment table, each payment row had its unique payment\_id ( a primary key) and identified the customer that made the payment through the customer\_id (a foreign key since it references the customer table's primary key)





- Example

Query Editor

Query History

1

SELECT \* FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- Primary Key for Payment Table

Query Editor

Query History

1

SELECT \* FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- Multiple Foreign Key References

Query Editor

Query History

1

SELECT \* FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- Note pgAdmin won't alert you to FK

Query Editor

Query History

1

SELECT \* FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- You may begin to realize primary key and foreign key typically make good column choices for joining together two or more tables.



- When creating tables and defining columns, we can use constraints to define columns as being a primary key, or attaching a foreign key relationship to another table.
- Let's quickly explore table properties in pgAdmin to see how to get information on primary and foreign keys!



# Constraints



- Constraints are the rules enforced on data columns on table.
- These are used to prevent invalid data from being entered into the database.
- This ensures the accuracy and reliability of the data in the database.





- Constraints can be divided into two main categories:
  - Column Constraints
    - Constrains the data in a column to adhere to certain conditions.
  - Table Constraints
    - applied to the entire table rather than to an individual column.



- The most common constraints used:
  - **NOT NULL** Constraint
    - Ensures that a column cannot have NULL value.
  - **UNIQUE** Constraint
    - Ensures that all values in a column are different.



- The most common constraints used:
  - **PRIMARY Key**
    - Uniquely identifies each row/record in a database table.
  - **FOREIGN Key**
    - Constrains data based on columns in other tables.



- The most common constraints used:
  - **CHECK** Constraint
    - Ensures that all values in a column satisfy certain conditions.



- The most common constraints used:
  - **EXCLUSION** Constraint
    - Ensures that if any two rows are compared on the specified column or expression using the specified operator, not all of these comparisons will return TRUE.



- Table Constraints
  - CHECK (condition)
    - to check a condition when inserting or updating data.
  - REFERENCES
    - to constrain the value stored in the column that must exist in a column in another table.



- Table Constraints
  - UNIQUE (column\_list)
    - forces the values stored in the columns listed inside the parentheses to be unique.
  - PRIMARY KEY(column\_list)
    - Allows you to define the primary key that consists of multiple columns.



- Now that we understand data types, primary keys, foreign keys, and constraints we are ready to begin using SQL syntax to create tables!





# CREATE



- Let's now learn the syntax to create a table in SQL using the CREATE keyword and column syntax.



- Full General Syntax
  - CREATE TABLE table\_name (  
column\_name TYPE column\_constraint,  
column\_name TYPE column\_constraint,  
table\_constraint table\_constraint  
) INHERITS existing\_table\_name;



- Full General Syntax
  - `CREATE TABLE table_name (  
column_name TYPE column_constraint,  
column_name TYPE column_constraint,  
table_constraint table_constraint  
) INHERITS existing_table_name;`



- Full General Syntax
  - CREATE TABLE table\_name (  
column\_name TYPE column\_constraint,  
column\_name TYPE column\_constraint,  
table\_constraint table\_constraint  
) INHERITS existing\_table\_name;



- Full General Syntax
  - CREATE TABLE table\_name (  
    column\_name TYPE column\_constraint,  
    column\_name TYPE column\_constraint,  
    table\_constraint table\_constraint  
) INHERITS existing\_table\_name;



- Full General Syntax
  - CREATE TABLE table\_name (  
column\_name TYPE column\_constraint,  
column\_name TYPE column\_constraint,  
table\_constraint table\_constraint  
) INHERITS existing\_table\_name;



- Full General Syntax
  - CREATE TABLE table\_name (  
column\_name TYPE column\_constraint,  
column\_name TYPE column\_constraint,  
table\_constraint table\_constraint  
) INHERITS existing\_table\_name;





- Full General Syntax

- CREATE TABLE table\_name (  
    column\_name TYPE column\_constraint,  
    column\_name TYPE column\_constraint,  
    table\_constraint table\_constraint  
) INHERITS existing\_table\_name;



- Common Simple Syntax
  - CREATE TABLE table\_name (  
    column\_name TYPE column\_constraint,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE **table\_name** (  
    column\_name TYPE column\_constraint,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE **players**(  
    column\_name TYPE column\_constraint,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE **players**(  
column\_name TYPE column\_constraint,  
column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE players(  
    column\_name TYPE column\_constraint,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE players(  
    player\_id TYPE column\_constraint,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE players(  
    player\_id **TYPE** column\_constraint,  
    column\_name TYPE column\_constraint,  
);





- Example Syntax
  - CREATE TABLE players(  
    player\_id **SERIAL** column\_constraint,  
    column\_name TYPE column\_constraint,  
);



- **SERIAL**
  - In PostgreSQL, a sequence is a special kind of database object that generates a sequence of integers.
  - A sequence is often used as the primary key column in a table.



- **SERIAL**
  - It will create a sequence object and set the next value generated by the sequence as the default value for the column.
  - This is perfect for a primary key, because it logs unique integer entries for you automatically upon insertion.



- **SERIAL**
  - If a row is later removed, the column with the SERIAL data type will **not** adjust, marking the fact that a row was removed from the sequence, for example
    - 1,2,3,5,6,7
      - You know row 4 was removed at some point



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



Name	Storage Size	Description	Range
smallint	2 bytes	small-range integer	-32768 to +32767
integer	4 bytes	typical choice for integer	-2147483648 to +2147483647
bigint	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
decimal	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
numeric	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
real	4 bytes	variable-precision, inexact	6 decimal digits precision
double precision	8 bytes	variable-precision, inexact	15 decimal digits precision
smallserial	2 bytes	small autoincrementing integer	1 to 32767
serial	4 bytes	autoincrementing integer	1 to 2147483647
bigserial	8 bytes	large autoincrementing integer	1 to 9223372036854775807



- Example Syntax
  - CREATE TABLE players(  
    player\_id **SERIAL** column\_constraint,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE players(  
    player\_id SERIAL **column\_constraint**,  
    column\_name TYPE column\_constraint,  
);





- Example Syntax
  - CREATE TABLE players(  
    player\_id SERIAL PRIMARY KEY,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE players(  
    player\_id SERIAL PRIMARY KEY,  
    column\_name TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE players(  
    player\_id SERIAL PRIMARY KEY,  
    age TYPE column\_constraint,  
);



- Example Syntax
  - CREATE TABLE players(  
player\_id SERIAL PRIMARY KEY,  
age **TYPE** column\_constraint,  
);



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



- Example Syntax
  - CREATE TABLE players(  
    player\_id SERIAL PRIMARY KEY,  
    age TYPE column\_constraint  
);



- Example Syntax
  - CREATE TABLE players(  
player\_id SERIAL PRIMARY KEY,  
age **SMALLINT** column\_constraint  
);





- Example Syntax
  - CREATE TABLE players(  
player\_id SERIAL PRIMARY KEY,  
age SMALLINT **column\_constraint**  
);



- Example Syntax
  - CREATE TABLE players(  
player\_id SERIAL PRIMARY KEY,  
age SMALLINT NOT NULL  
);



- Example Syntax
  - CREATE TABLE players(  
    player\_id SERIAL PRIMARY KEY,  
    age SMALLINT NOT NULL  
);



- Let's explore some examples in pgAdmin!



# INSERT



- INSERT allows you to add in rows to a table.
- General Syntax
  - INSERT INTO table (column1, column2, ...) VALUES  
    (value1, value2, ...),  
    (value1, value2, ...) ,...;



- INSERT allows you to add in rows to a table.
- Syntax for Inserting Values from another table:
  - `INSERT INTO table(column1,column2,...)`  
`SELECT column1,column2,...`  
`FROM another_table`  
`WHERE condition;`



- Keep in mind, the inserted row values must match up for the table, including constraints.
- SERIAL columns do not need to be provided a value.
- Let's use INSERT in pgAdmin!





# UPDATE



- The UPDATE keyword allows for the changing of values of the columns in a table.



- General Syntax
  - UPDATE table  
SET column1 = value1,  
column2 = value2 ,...  
WHERE  
condition;



- Example
  - UPDATE account  
SET last\_login = CURRENT\_TIMESTAMP  
WHERE last\_login IS NULL;



- Reset everything without WHERE condition
  - UPDATE account  
SET last\_login = CURRENT\_TIMESTAMP



- Set based on another column
  - UPDATE account  
SET last\_login = created\_on



- Using another table's values (UPDATE join)
  - UPDATE TableA  
SET original\_col = TableB.new\_col  
FROM tableB  
WHERE tableA.id = TableB.id



- Return affected rows
  - UPDATE account  
SET last\_login = created\_on  
RETURNING account\_id,last\_login





- Let's explore this further in pgAdmin!



# DELETE



- We can use the DELETE clause to remove rows from a table.
- For example:
  - DELETE FROM table  
WHERE row\_id = 1



- We can delete rows based on their presence in other tables
- For example:
  - `DELETE FROM tableA  
USING tableB  
WHERE tableA.id=TableB.id`



- We can delete all rows from a table
- For example:
  - DELETE FROM table



- Similar to UPDATE command, you can also add in a RETURNING call to return rows that were removed.
- Let's explore DELETE with pgAdmin!



# ALTER



- The ALTER clause allows for changes to an existing table structure, such as:
  - Adding, dropping, or renaming columns
  - Changing a column's data type
  - Set DEFAULT values for a column
  - Add CHECK constraints
  - Rename table





- General Syntax
  - ALTER TABLE table\_name action



- Adding Columns
  - ALTER TABLE table\_name  
ADD COLUMN new\_col TYPE



- Removing Columns
  - ALTER TABLE table\_name  
DROP COLUMN col\_name



- Alter constraints
  - ALTER TABLE table\_name  
ALTER COLUMN col\_name  
SET DEFAULT value



- Alter constraints
  - ALTER TABLE table\_name
  - ALTER COLUMN col\_name
  - DROP DEFAULT



- Alter constraints
  - ALTER TABLE table\_name  
ALTER COLUMN col\_name  
SET NOT NULL



- Alter constraints
  - ALTER TABLE table\_name  
ALTER COLUMN col\_name  
DROP NOT NULL



- Alter constraints

- ALTER TABLE table\_name

- ALTER COLUMN col\_name

- ADD CONSTRAINT constraint\_name





- Let's explore some examples in pgAdmin!



# DROP



- DROP allows for the complete removal of a column in a table.
- In PostgreSQL this will also automatically remove all of its indexes and constraints involving the column.
- However, it will not remove columns used in views, triggers, or stored procedures without the additional CASCADE clause.



- General Syntax
  - ALTER TABLE table\_name
  - DROP COLUMN col\_name



- Remove all dependencies
  - ALTER TABLE table\_name

DROP COLUMN col\_name CASCADE



- Check for existence to avoid error
  - ALTER TABLE table\_name

DROP COLUMN IF EXISTS col\_name



- Drop multiple columns
  - ALTER TABLE table\_name

DROP COLUMN col\_one,

DROP COLUMN col\_two



- Let's see a quick example in pgAdmin!





# CHECK



- The CHECK constraint allows us to create more customized constraints that adhere to a certain condition.
- Such as making sure all inserted integer values fall below a certain threshold.



- General Syntax
  - CREATE TABLE example(  
ex\_id SERIAL PRIMARY KEY,  
age SMALLINT CHECK (age > 21),  
parent\_age SMALLINT CHECK (  
parent\_age > age)  
);



- Let's explore this concept in pgAdmin!